

Implémentation du multicœur en couche logicielle basse

Summary

- Objectif
- Qu'est-ce qu'un Spinlock ?
- Monocœur → Multicœur
- Mise en place d'un big lock
- Passage aux small spinlocks
- Conclusion

Objectif

**Support multicœur et gestion de
la concurrence**

Concurrence

Cœur 0 :

uart_send_string("Hello, World!")

Cœur 2 :

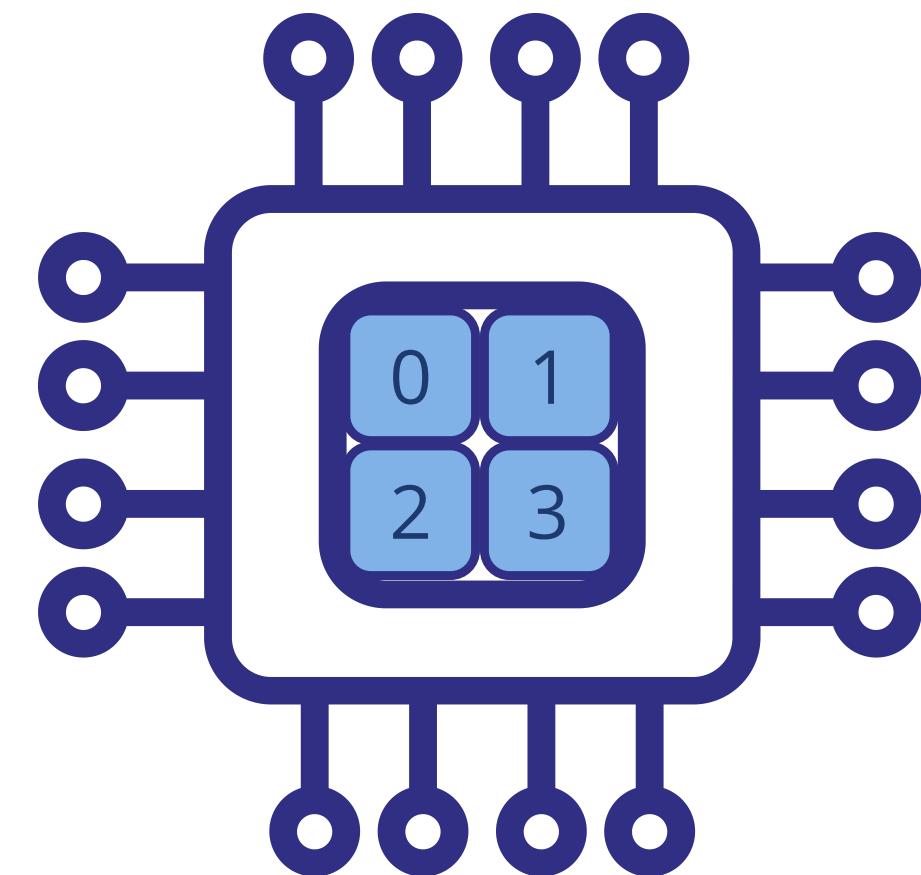
uart_send_string("Hello, World!")

Cœur 1 :

uart_send_string("Hello, World!")

Cœur 3 :

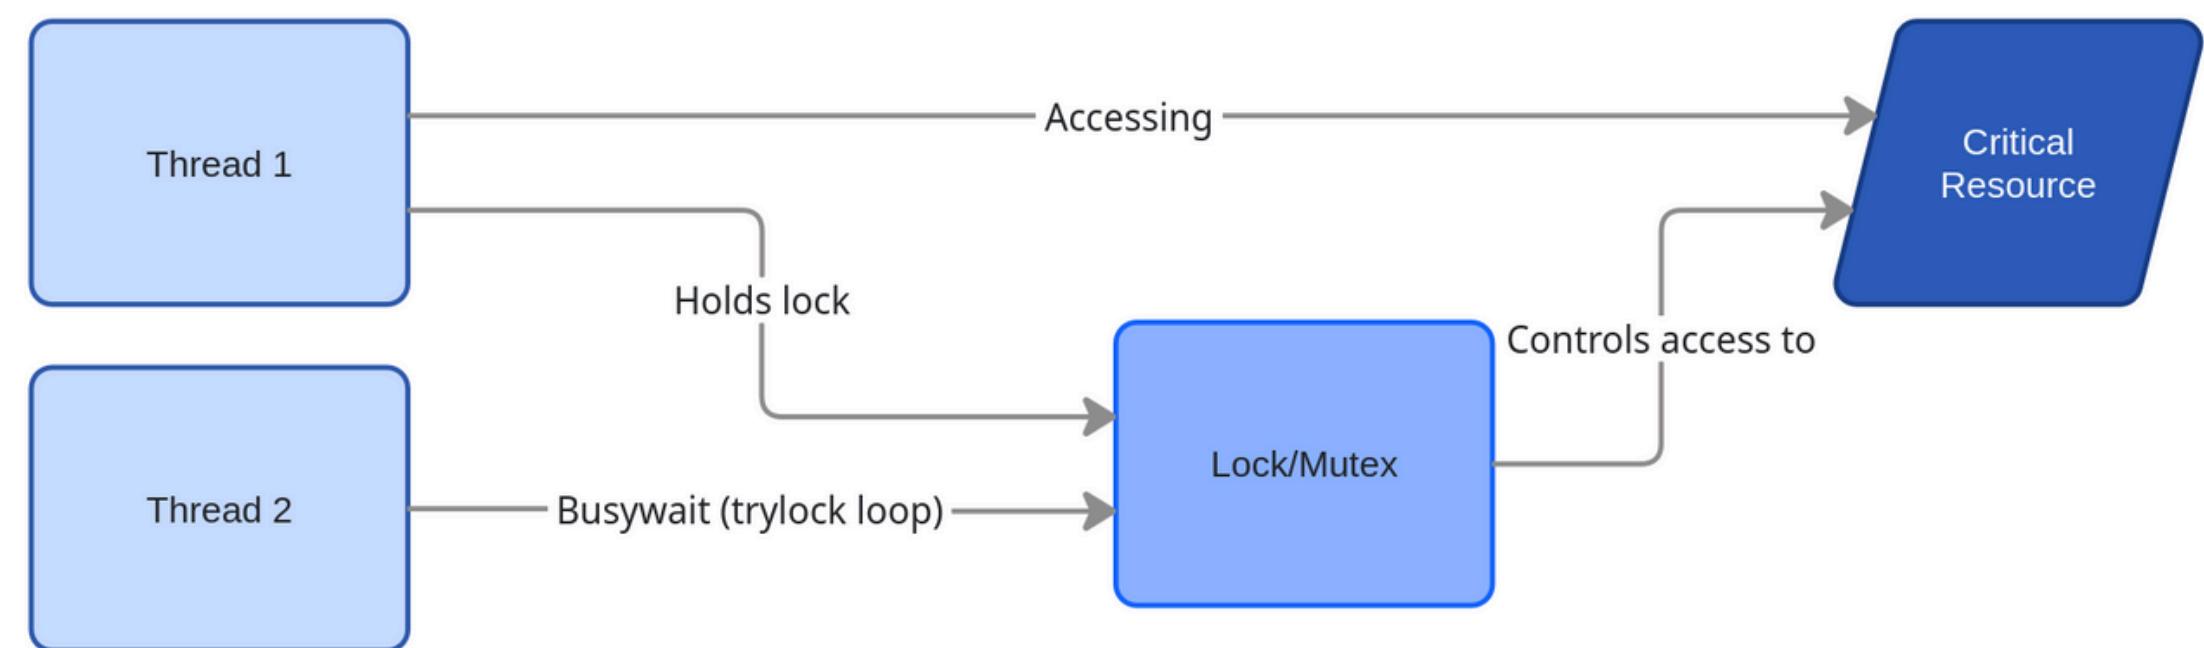
uart_send_string("Hello, World!")



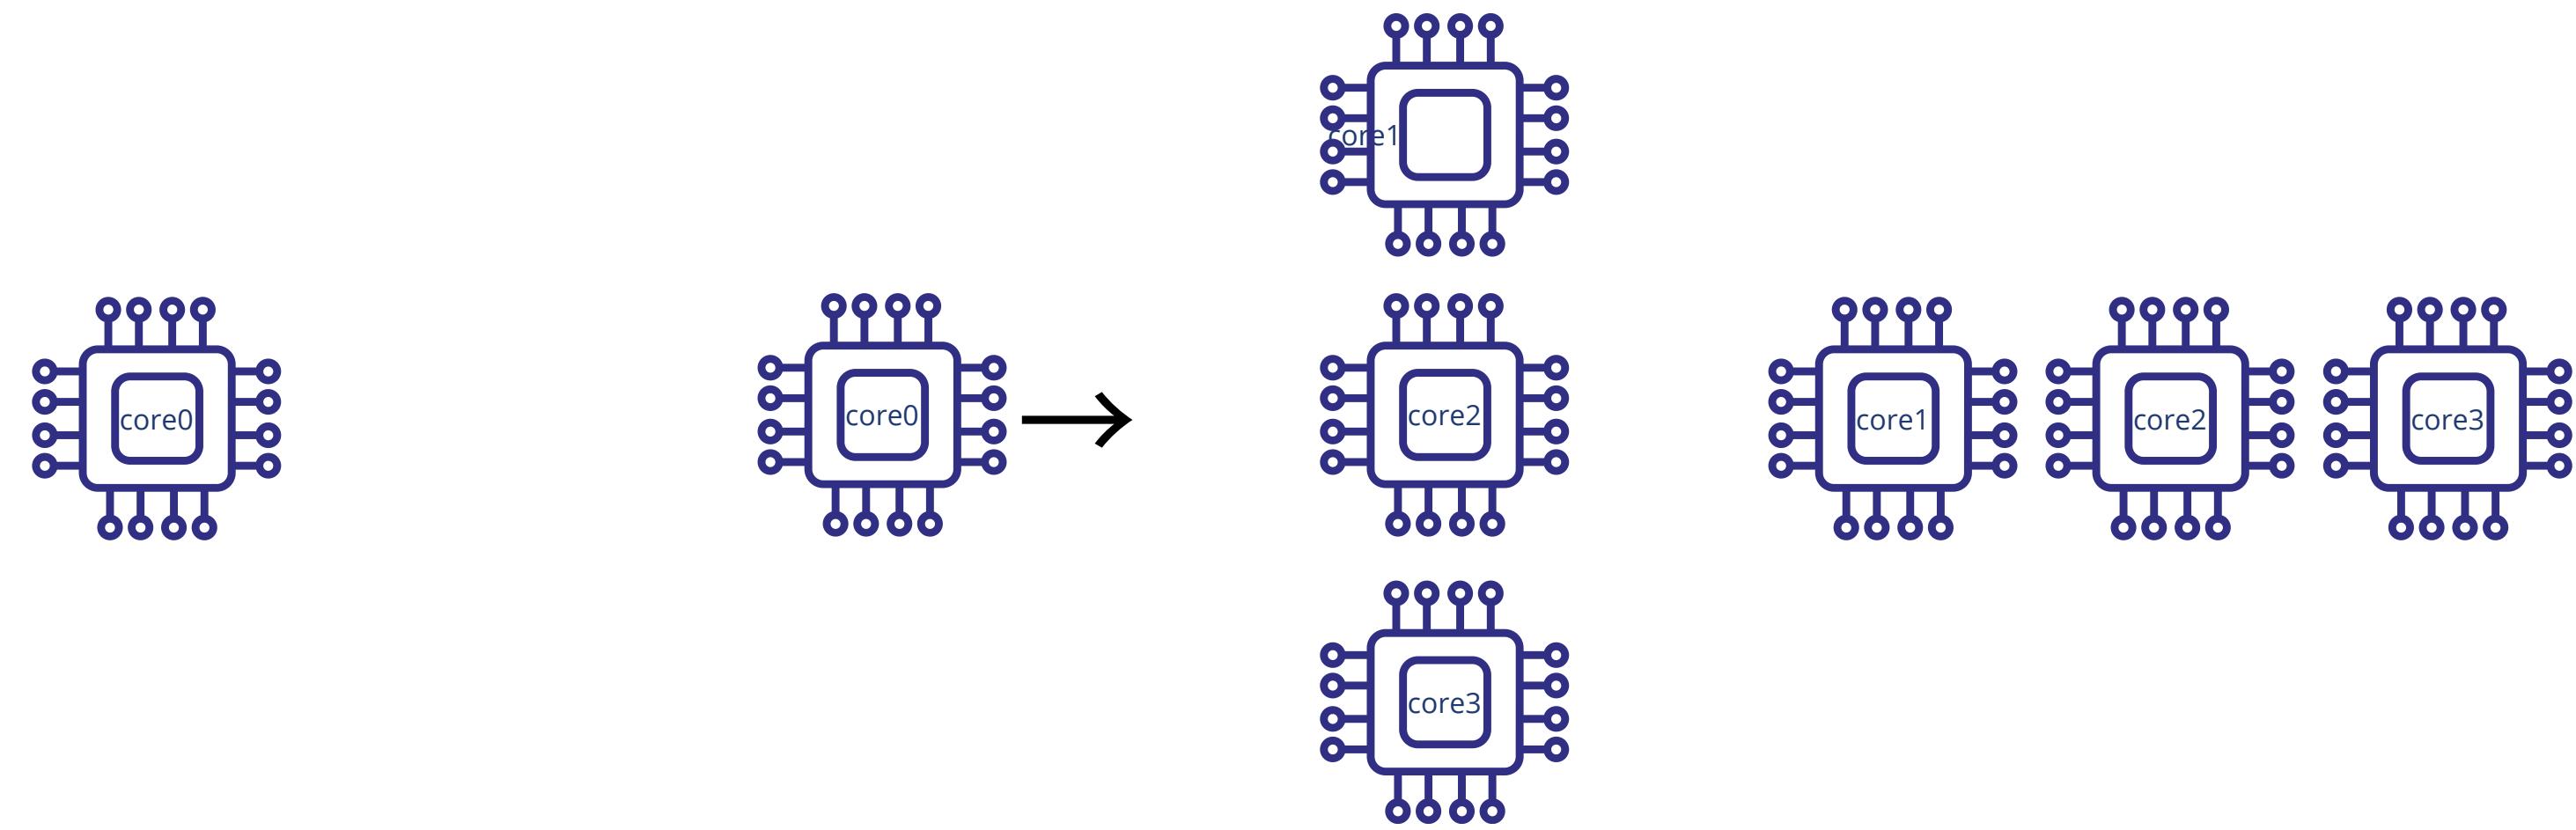
Sortie :

HHeelllloo,, WWHoerrllldo!!, HWeolrllod,! World!

Qu'est-ce qu'un spinlock ?



Monocœur → Multicoeur



Etape 1 : Initialiser notre système et le cœur 0

Etape 2 : Réveiller les autres cœurs

Etape 3 : Initialiser tous les cœurs

Mono cœur → Multicœur

```
.globl start_cores
start_cores:
    adrp x2, _start
    add x2, x2, #:lo12:_start
    ldr x3, =VA_START
    sub x2, x2, x3 // Since t
give the other cores a phys:
virtual offset
    mov x1, #0xd8
    str x2, [x1,#8]!
    str x2, [x1,#8]!
    str x2, [x1,#8]
    ret
```

Etape 2 : Réveiller les autres coeurs

```
if (core_id == 0){
    uart_init();
    init_printf(NULL, putc);
    printf("kernel boots ...\\n\\r");
    enable_interrupt_controller();
    int res = copy_process(PF_KTHREAD, (unsigned
long)&kernel_process, 0);
    if (res < 0) {
        printf("error while starting kernel process");
        return;
    }
    start_cores();
    printf("Starting other cores...\\n\\r");
}
irq_vector_init();
generic_timer_init();
enable_irq();

printf("Core %d : Successfully started\\n", core_id);
unlock();
while(1){
    lock();
    schedule(core_id);
    unlock();
    asm("wfi");
}
```

Etape 1 : Initialiser notre système et le cœur 0

```
// Fetch PMD address
adrp x11, pg_dir
add x11, x11, #(3 * PAGE_SIZE)

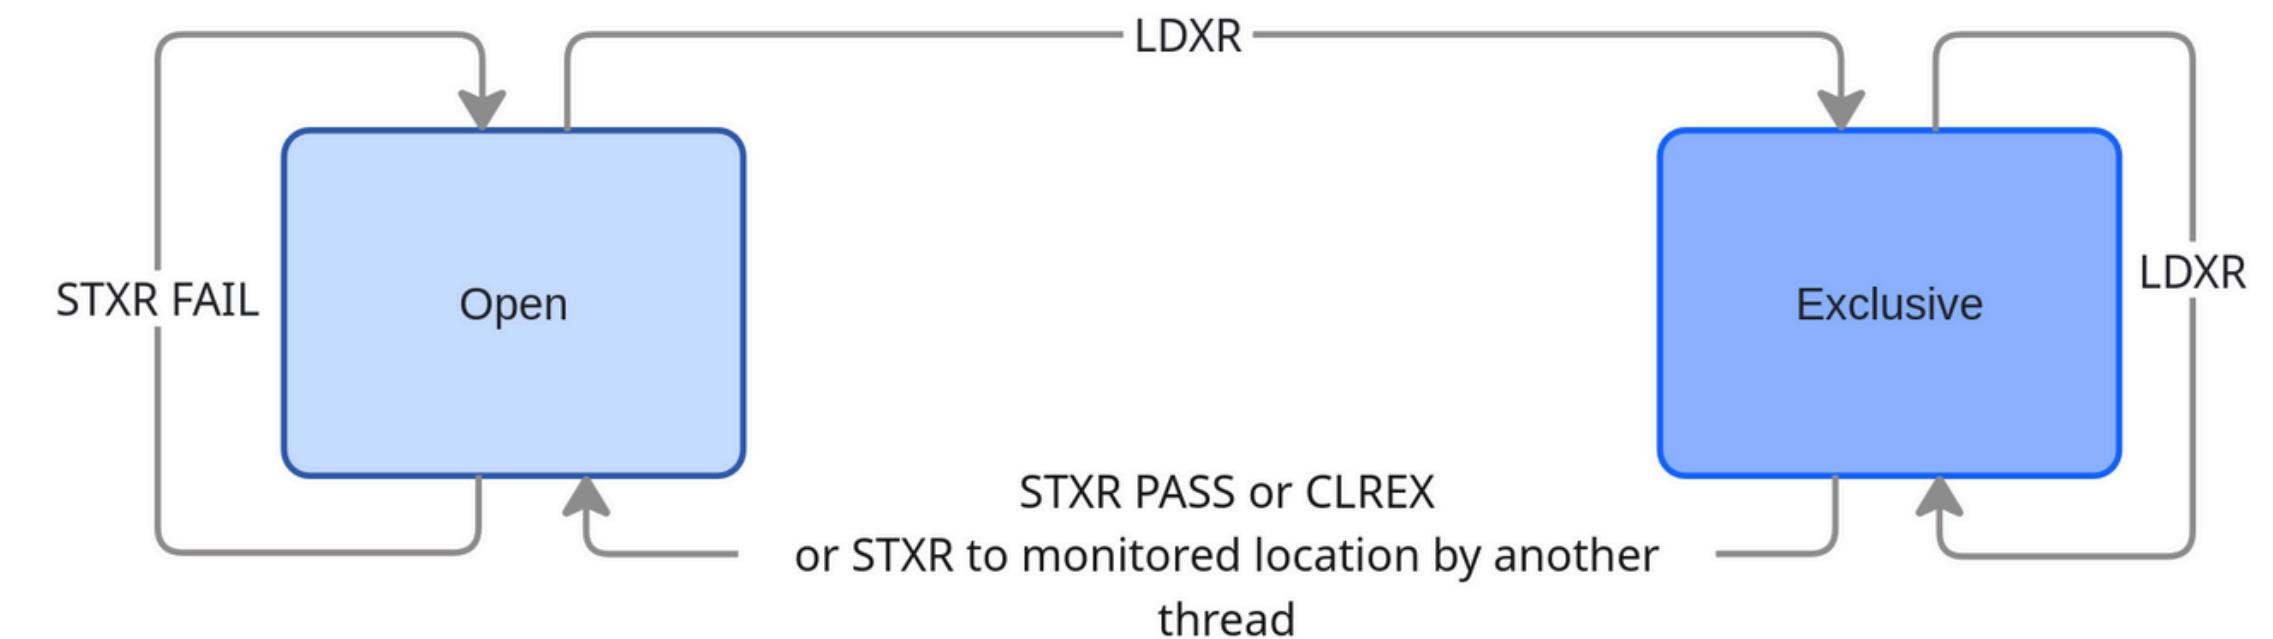
// Map second PMD address
orr x12, x11, #MM_TYPE_PAGE_TABLE
str x12, [x10, #8]

// Map timer addresses
mov x0, x11
ldr x1, =GENERIC_TIMER_START // Physical
address of timer region
ldr x2, =(VA_START + GENERIC_TIMER_START) // Virtual
address start
ldr x3, =(VA_START + GENERIC_TIMER_START) // Virtual
address end : same address as start to enforce the mapping
of a single section

void enable_interrupt_controller()
{
    // Enables Core 0-3 Timers interrupt control for the
generic timer
    put32(TIMER_INT_CTRL_0, TIMER_INT_CTRL_0_VALUE);
    put32(TIMER_INT_CTRL_1, TIMER_INT_CTRL_0_VALUE);
    put32(TIMER_INT_CTRL_2, TIMER_INT_CTRL_0_VALUE);
    put32(TIMER_INT_CTRL_3, TIMER_INT_CTRL_0_VALUE);
}
```

Etape 3 : Initialiser tous les coeurs

Implémentation: Spinlock



Load exclusif et Store exclusif

- ldxrb : enregistre l'adresse physique en tant qu'accès exclusif au cœur qui a load
- stxrb : store dans l'adresse physique à condition qu'elle soit enregistrée

```
.globl lock
lock:
    mrs x9, mpidr_el1
    and x9, x9, #0xFF           // x9 = core_id
    adrp x10, mutex
    add x10, x10, #:lo12:mutex

    // Check ownership
    ldrb w11, [x10]
    cmp w11, w9
    b.eq 1f

    // Spinloop
spinloop:
    ldxrb w11, [x10]
    cmp w11, #0xFF
    b.ne spinloop // Check mutex availability
    stxrb w11, w9, [x10]
    cbnz w11, spinloop // Check if memory access was exclus
    dmb sy

1: ret
```

```
.globl unlock
unlock:
    mrs x9, mpidr_el1
    and x9, x9, #0xFF
    adrp x10, mutex
    add x10, x10, #:lo12:mutex

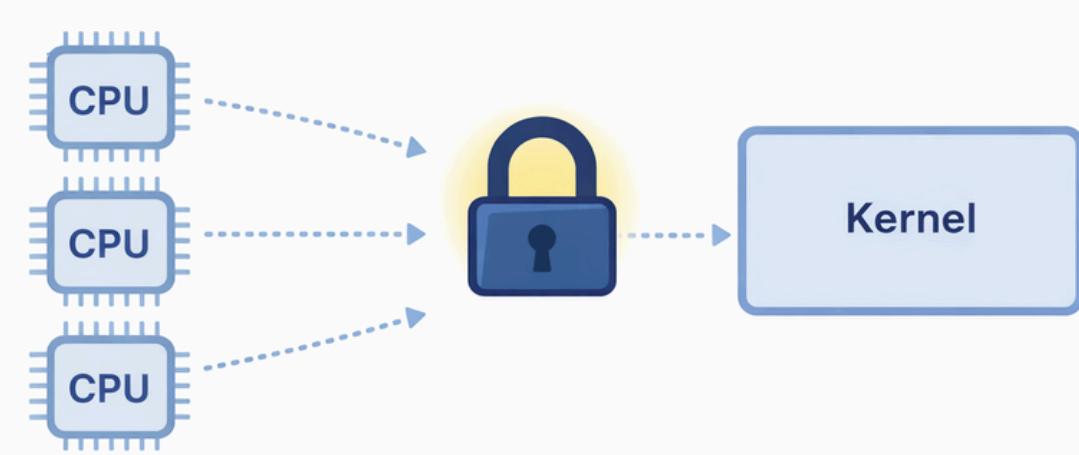
    ldrb w11, [x10]
    cmp w11, w9
    b.ne 1f

    dmb sy
    mov w11, #0xFF
    strb w11, [x10]

1: ret

.section ".data"
.globl mutex
mutex :
    .byte 0xFF
```

Le Big Lock

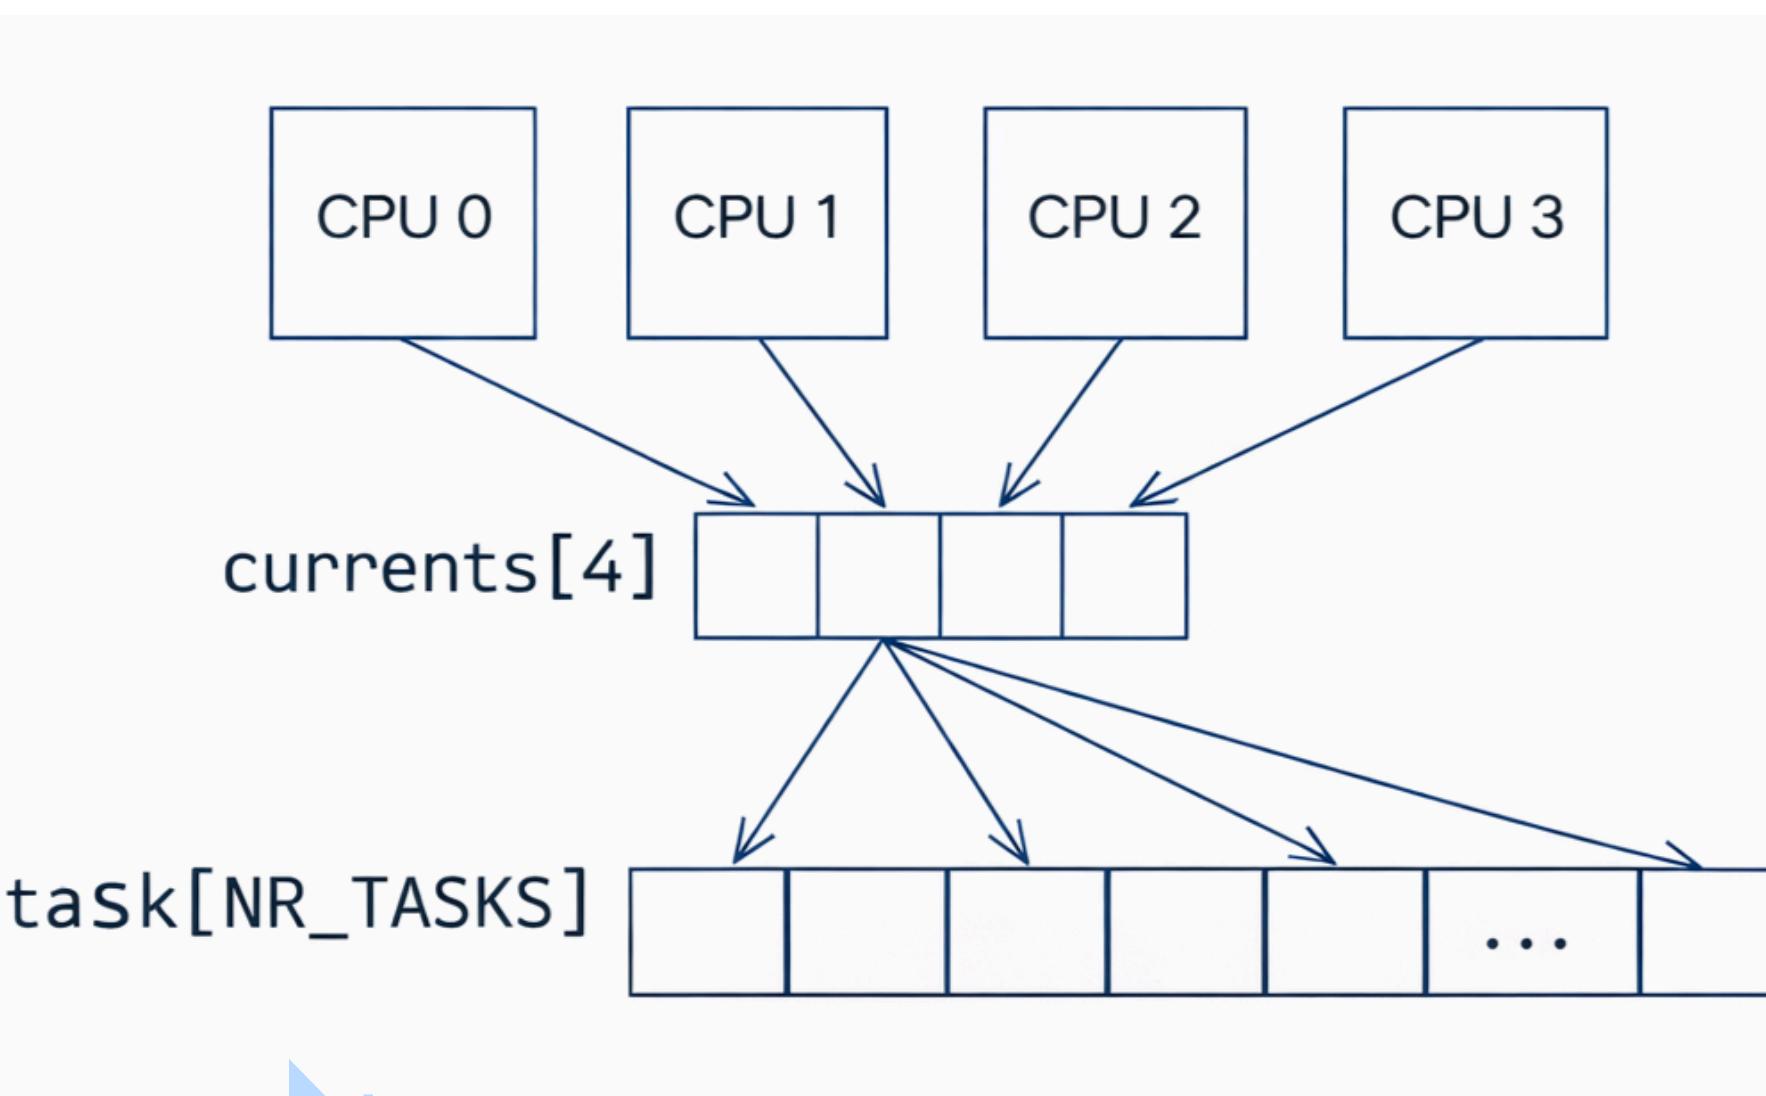


entry.S

```
ell_irq:  
    kernel_entry 1  
    bl lock // KERNEL LOCK  
    bl handle_irq  
    bl unlock // KERNEL UNLOCK  
    kernel_exit 1  
  
el0_irq:  
    kernel_entry 0  
    bl lock // KERNEL LOCK  
    bl handle_irq  
    bl unlock // KERNEL UNLOCK  
    kernel_exit 0  
  
el0_sync:  
    kernel_entry 0  
    bl lock // KERNEL LOCK  
    ...  
  
el0_svc:  
    ...  
  
ret_from_syscall:  
    str x0, [sp, #S_X0]      // returned x0  
    bl unlock // KERNEL UNLOCK  
    kernel_exit 0
```

```
el0_da:  
    ...  
    bl unlock // KERNEL UNLOCK  
    kernel_exit 0  
  
.globl ret_from_fork  
ret_from_fork:  
    bl schedule_tail  
    bl unlock // KERNEL UNLOCK  
    cbz x19, ret_to_user      // not a kernel thread  
    mov x0, x20  
    blr x19  
ret_to_user:  
// bl unlock // Free the lock applied to kernel process  
    bl disable_irq  
    kernel_exit 0
```

Nouveau Scheduler

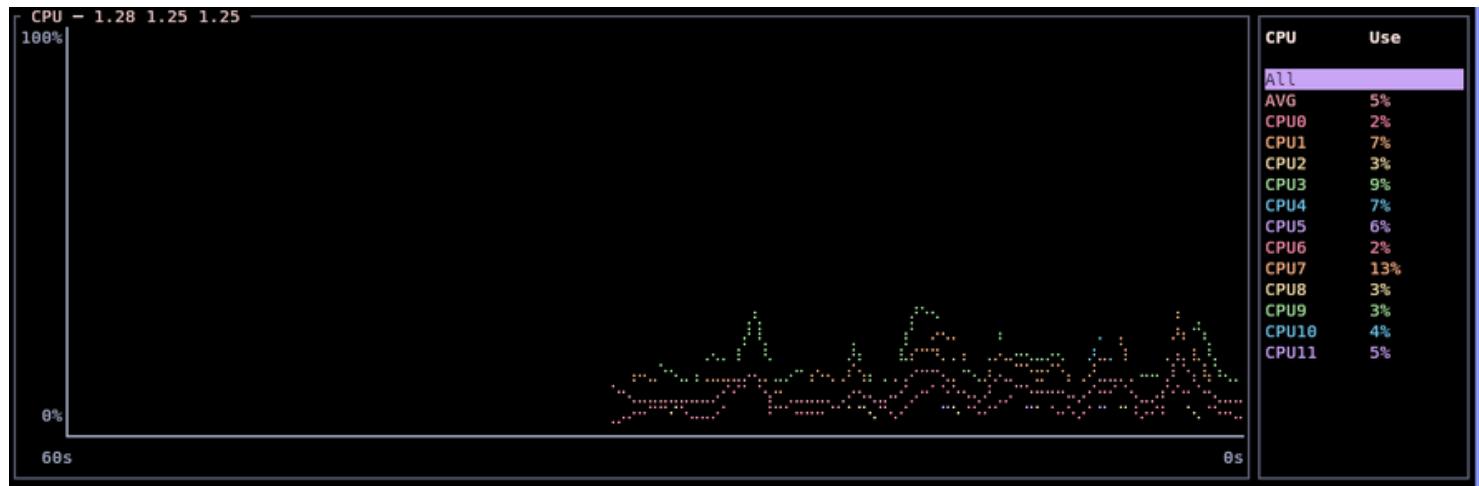


```
static struct task_struct init_task_0 = INIT_TASK;
static struct task_struct init_task_1 = INIT_TASK;
static struct task_struct init_task_2 = INIT_TASK;
static struct task_struct init_task_3 = INIT_TASK;
struct task_struct *currents[NB_CPU] = { &(init_task_0), &(init_task_1),
&(init_task_2), &(init_task_3)};
struct task_struct *task[NR_TASKS] = {
    &(init_task_0),
    &(init_task_1),
    &(init_task_2),
    &(init_task_3),
};

int nr_tasks = NB_CPU;

struct task_struct {
    struct cpu_context cpu_context;
    long state;
    long counter;
    long priority;
    long preempt_count;
    unsigned long flags;
    struct mm_struct mm;
    unsigned char taken;
};
```

Méthode de test de la v1



- Bottom

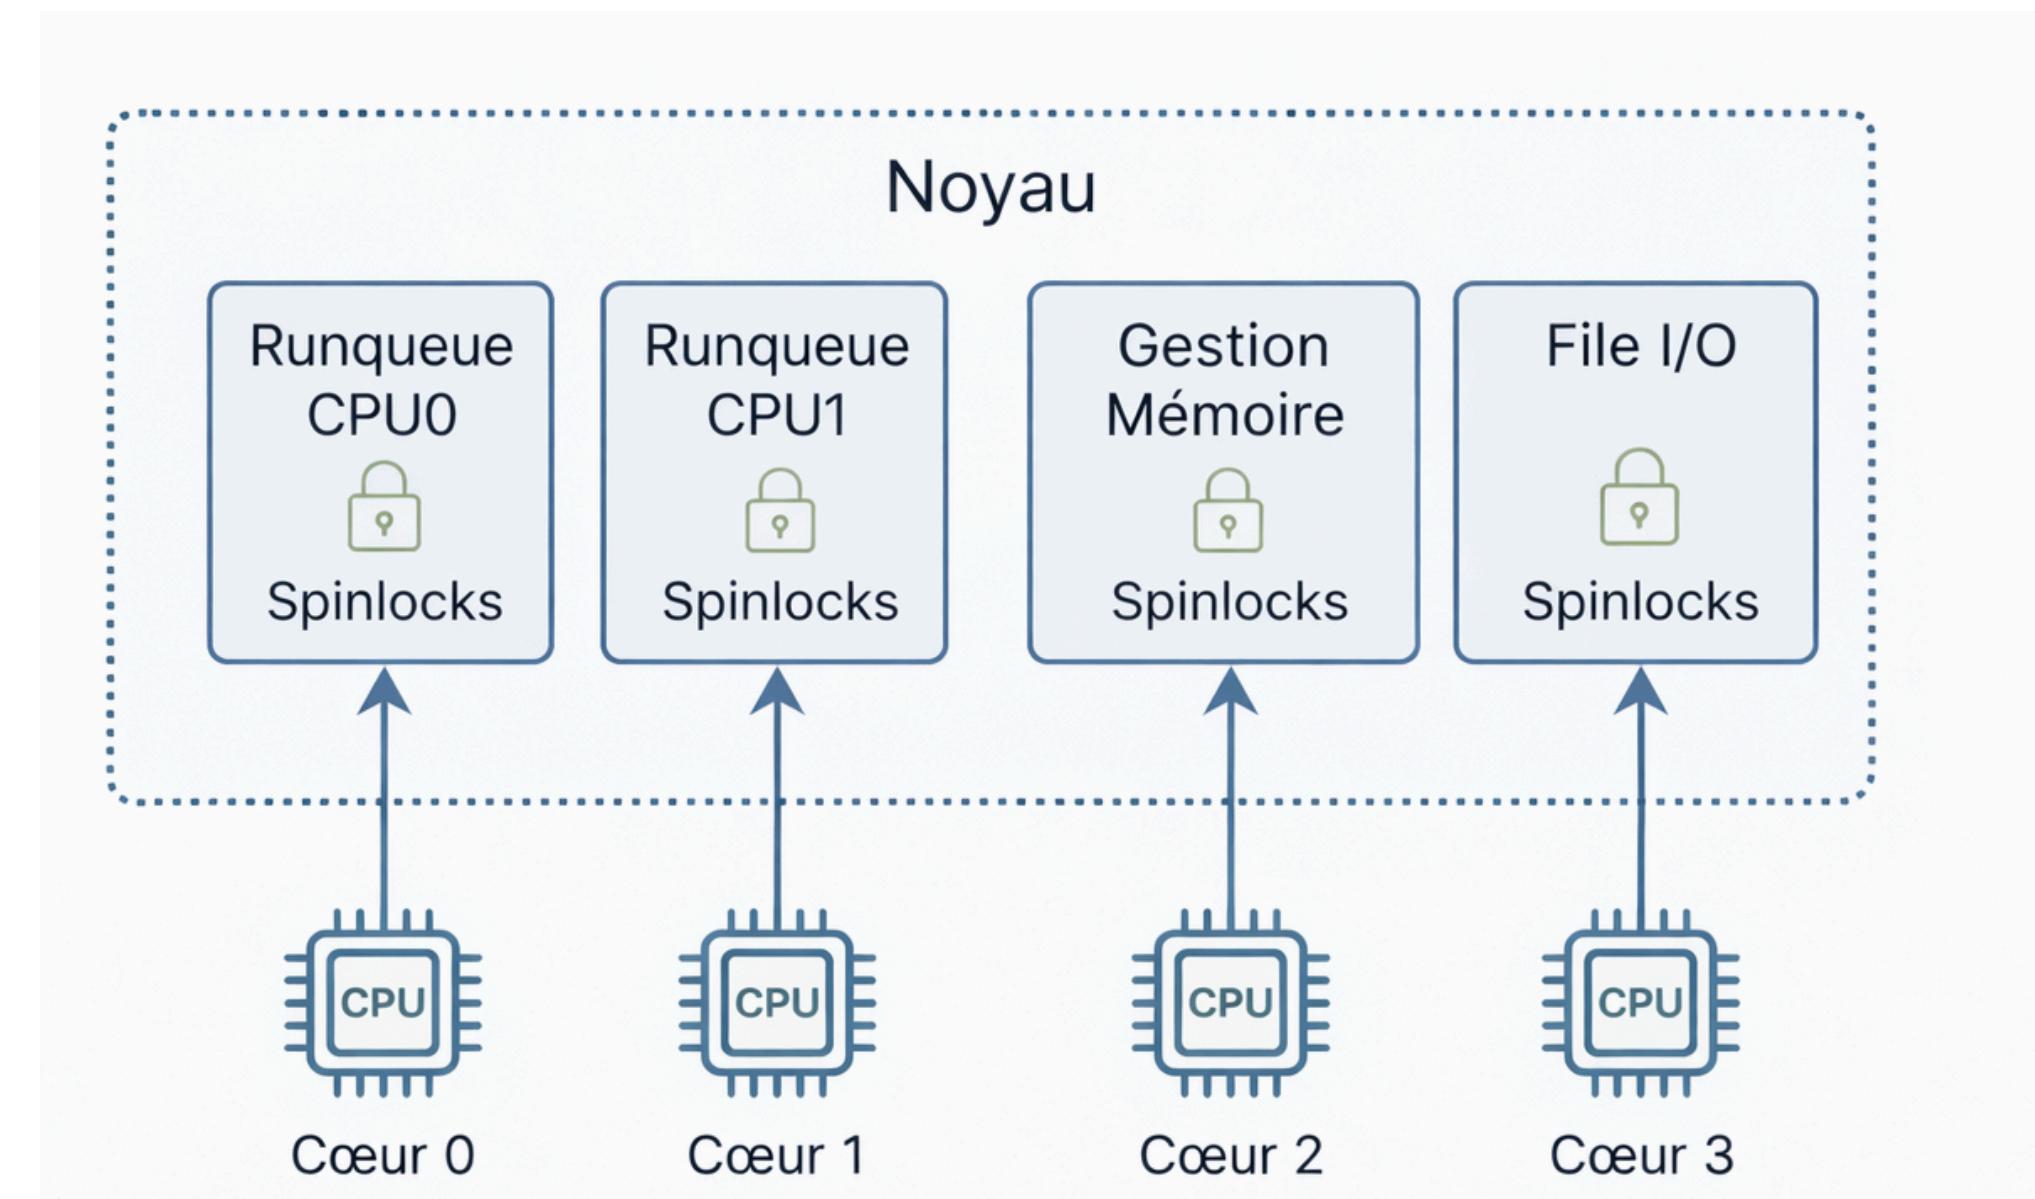
```
void create_and_loop(char* str, char* arg)
{
    char buf[25] = {"\0"};
    for (int i = 0; i < 25; i++){
        buf[i] = str[i];
    }
    call_sys_write(buf);
    int pid = call_sys_fork();
    if (pid < 0) {
        call_sys_write("Error during fork\n\r");
        call_sys_exit();
        return;
    }
    if (pid == 0)
        loop(arg);

    call_sys_write("Process over\n\r");
    call_sys_exit();
}
```

```
void user_process()
{
    call_sys_write("User process\n\r");
    // First child
    call_sys_write("Creating process 1\n\r");
    int pid = call_sys_fork();
    if (pid < 0) {
        call_sys_write("Error during fork\n\r");
        call_sys_exit();
        return;
    }
    if (pid == 0)
        create_and_loop("Child 1 creating
child\n\r","azazazaz");

    // Second child ..
    // Third child..
    // Four child ..
}
```

Principe des Small Spinlocks



Implémentation des Small Spinlocks

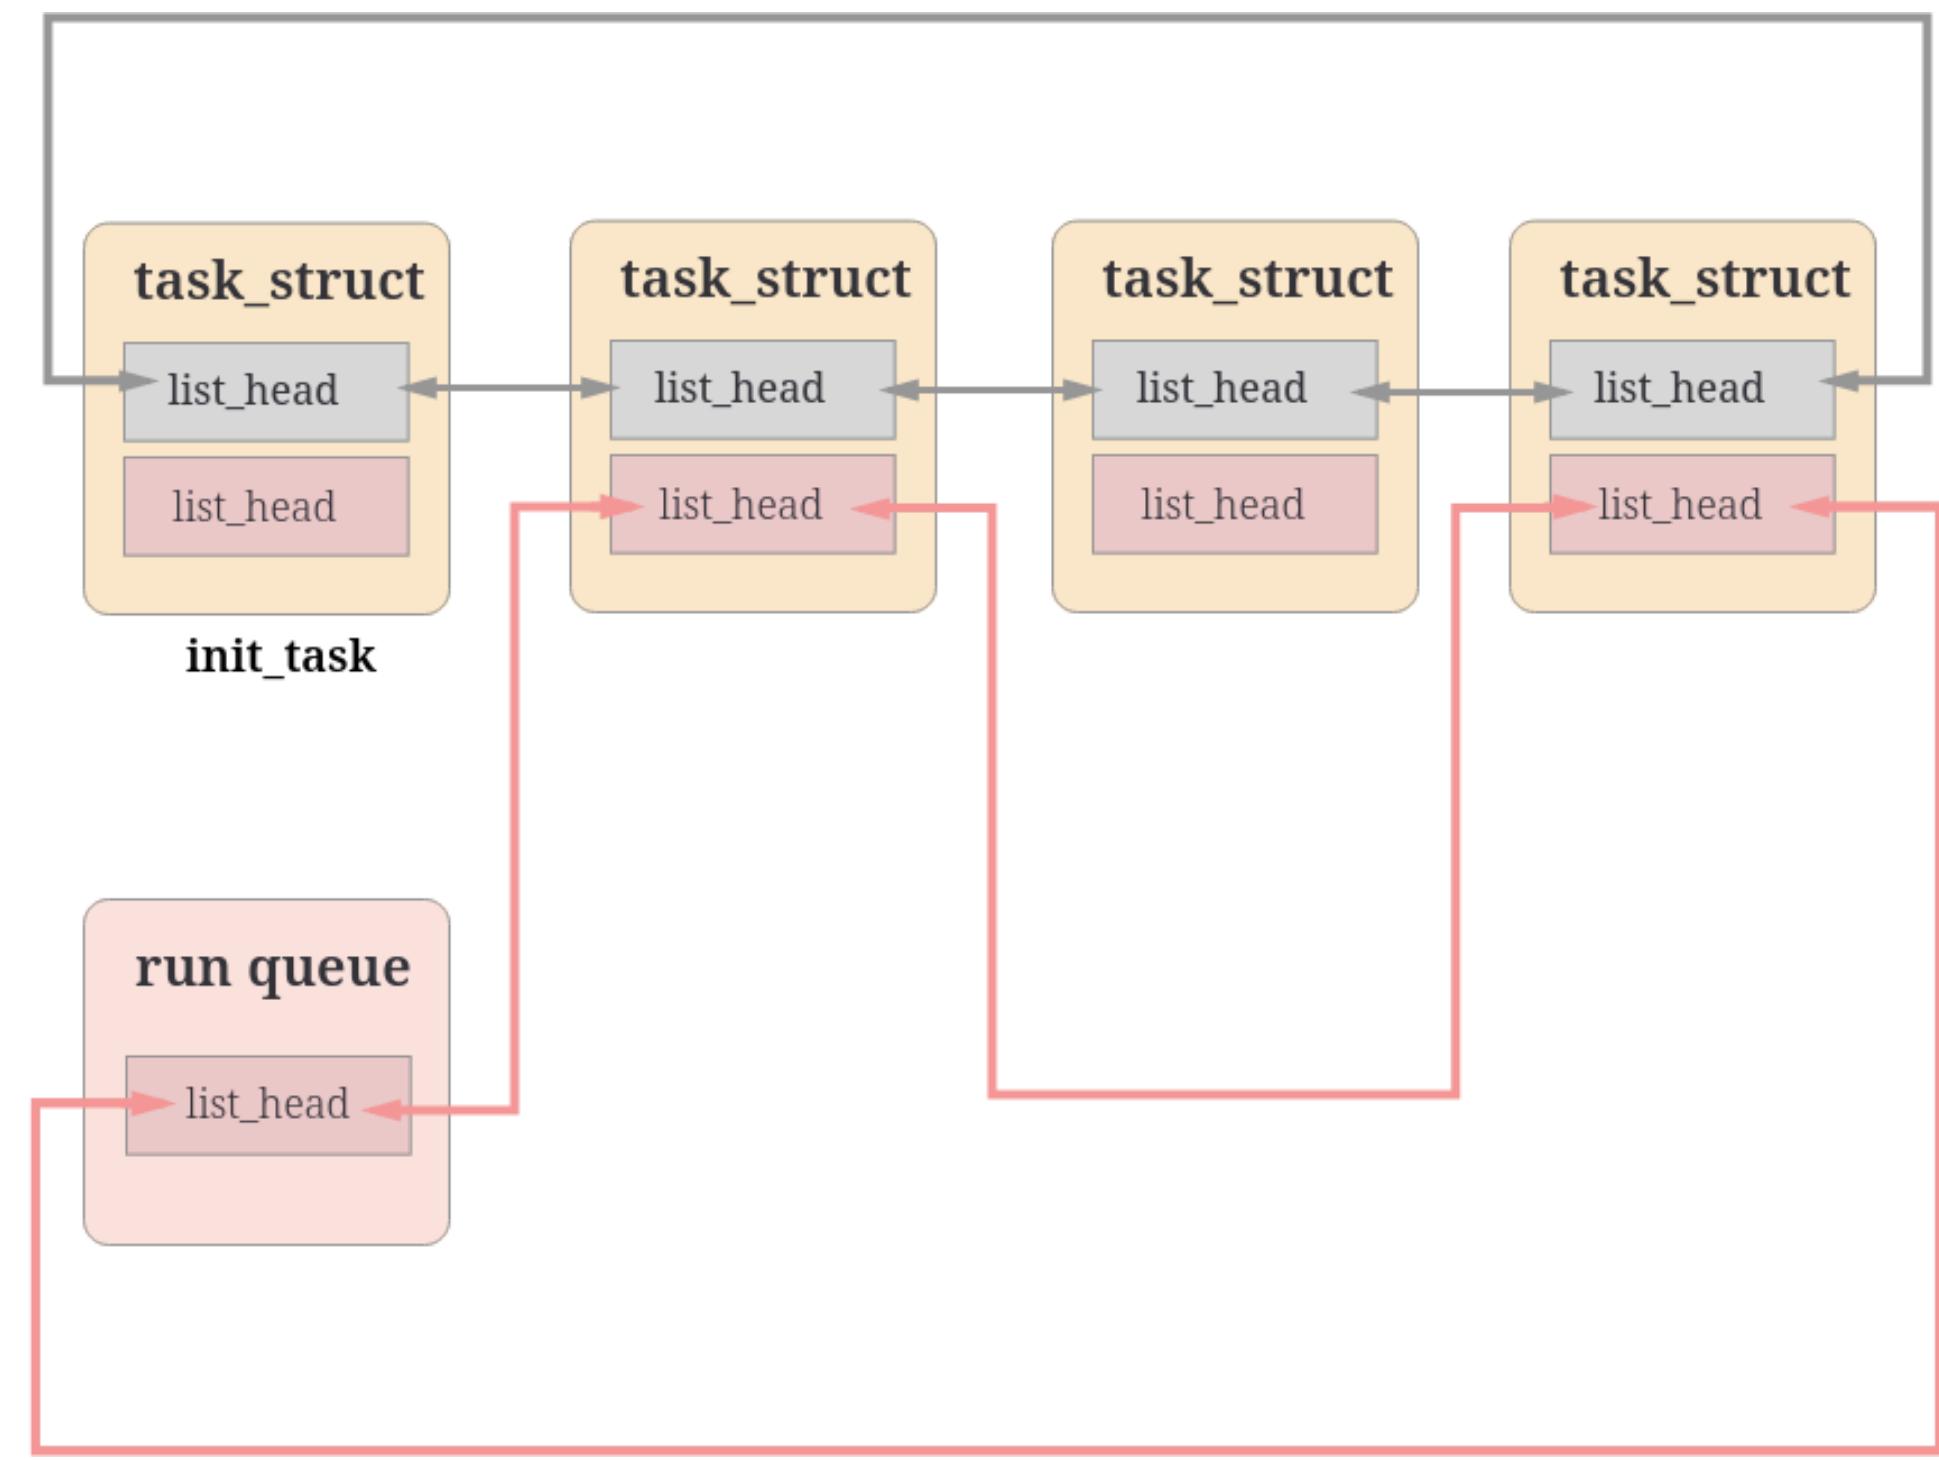
Nouvelles variables

```
.section ".data"
.globl mutex_write
mutex_write :
    .byte 0xFF
.globl mutex_shed
mutex_sched :
    .byte 0xFF
.globl mutex_mem
mutex_mem :
    .byte 0xFF
```

Addition atomique

```
.globl atomic_add
atomic_add:
    ldxr w2, [x0]
    add w2, w2, w1
    stxr w3, w2, [x0]
    cbnz w3, atomic_add
    dmb sy
    mov w0, w2
    ret
```

Optimisation : Scheduling per-cpu



Conclusion

- Alternatives:
 - CAS
 - WFE
- Axes d'amélioration:
 - équilibrage de charge (load balancing)
 - migration dynamique des tâches entre cœurs